

Classification and Regression

From linear and logistic regression to neural networks

FYS-STK4155: Applied data analysis and machine learning
Department of Physics, University of Oslo

Mina Tangen and Tham Le

 <https://github.com/minatang1/project2>

November, 2019

Abstract

We tried to implement logistic regression and our own neural network for classification. The data used is credit card default data, and the aim was to use neural networks to predict whether a customer defaults or not. The neural network was changed to do regression on Franke's function. The results obtained from this was compared to OLS, Ridge and LASSO regression. Our own implementation was compared to that of scikit-learn. Evaluation metrics such as F1 score, R2 and MSE was evaluated. We found that neural networks performs best in both cases. Scikit-learn performs, in general, much better than our implementation.

Contents

1	Introduction	3
2	Methods	4
2.1	Logistic Regression	4
2.1.1	Binary classification and the Sigmoid function	4
2.1.2	Maximum likelihood and the cost function	4
2.2	Stochastic Gradient Descent Methods (SGD)	5
2.3	Multilayer preceptron (MLP)	6
2.3.1	Neurons	6
2.3.2	Activation function	7
2.3.3	Feedforward and backpropagation	9
2.4	One hot encoding	10
2.5	Model evaluation	10
3	Data	12
3.1	Credit card data	12
3.2	Franke function	13
4	Implementation	13
4.1	Data preprocessing	13
4.2	Evaluation metrics	15
5	Results & discussion	18
5.1	Neural network	18
5.2	Franke's function	26
6	Conclusion	31
7	References	32
8	Appendix	34
8.1	Explanatory variables	34

1 Introduction

Classification in statistical analysis is an important tool to identifying the outcomes in different situations and sort large amount of datasets. We want to know if the outcomes are predicted and classified right. Examples of classification problems includes to classify a given physical state as solid or not, a pasient as healthy or sick or an e-mail as spam or not. In this project we study the challenges associated with classification and regression through implementation of logistic regression and the multilayer precepton (MLP).

The dataset we use for this project is retrieved from UCI Machine Learning Repository and contains default payments from a bank in Taiwan (Yeh, I-C. & Lien, C-H. (2009)). Our aim using this data set is to train a network, so that the explanatory variables X_i are used to predict the response variable Y . Or, stated in more colloquial terms, we want to train our network to identify how a client's gender, age, education, previous payment history etc. can explain whether or not he/she pays their credit card bill on time.

In the next section we will review methods such as logistic regression, stochastic gradient descent and multilayer precepton and their applications on the data set (Yeh, I-C. & Lien, C-H. (2009)). Iin the following sections the results will be presented and discussed. Then, we will reach a conclusion in the last section. Additional equations, information about the dataset and code will be presented in the appendix.

2 Methods

2.1 Logistic Regression

2.1.1 Binary classification and the Sigmoid function

The logistic model classifies the observed data x_i to be binary outcomes such as passed/failed, winning/losing, spam/not spam etc. Classification aims to predict the outcome correctly by finding patterns based on discrete variables. The probability that x_i belongs to a category $y_i = (0, 1)$ is given by the sigmoid function (Hjort-Jensen, M. (2019))

$$f(x) = \frac{1}{1 + \exp^{-x}} \quad (1)$$

and the probability of these two classes with y_i either 0 or 1.

$$p(y_i = 1|x_i, \hat{\beta}) = \frac{\exp(\beta_0 + \beta_1 x_i)}{1 + \exp(\beta_0 + \beta_1 x_i)}$$
$$p(y_i = 0|x_i, \hat{\beta}) = 1 - p(y_i = 1|x_i, \hat{\beta})$$

where $\hat{\beta}$ are the weights.

2.1.2 Maximum likelihood and the cost function

To find the total likelihood for all possible outcomes from the credit card dataset $D = y_i, x_i$, with the binary labels $y_i = (0, 1)$ we can use the Maximum Likelihood Estimation (MLE). MLE is the product of the individual probabilities for every outcome y_i (Hjort-Jensen, M. (2019)) defined as

$$P(\mathcal{D}|\hat{\beta}) = \prod_{i=1}^n [p(y_i = 1|x_i, \hat{\beta})]^{y_i} [1 - p(y_i = 1|x_i, \hat{\beta})]^{1-y_i} \quad (2)$$

and from the negative log-likelihood function, we can obtain the cost function as

$$-P(\mathcal{D}|\hat{\beta})_{\log} = C(\hat{\beta}) = -\sum_{i=1}^n y_i(\beta_0 + \beta_1 x_i) - \log(1 + \exp \beta_0 + \beta_1 x_i) \quad (3)$$

2.2 Stochastic Gradient Descent Methods (SGD)

To find the best possible fit for the dataset, we need to optimize the cost function $C(\hat{\beta})$ by finding its minimum (Hjort-Jensen, M. (2019)). The SGD method is commonly used to achieve this. The general idea is that a function $F(x)$ decreases from x in a negative direction with negative gradient $-\nabla F(x_i)$ where x are iteratively adjusted. The function is expressed as

$$x_{i+1} = x_i - \eta \nabla F(x_i) \quad (4)$$

with $\eta > 0$. The parameter η is the learning rate, also the step length. It is important to choose a reasonable value for η , because small values will take longer time for SGD to converge or maybe not converge at all within the given time frame. But for larger values SGD might pass the minimum or be too unstable to calculate.




Figure 1: Gradient descent algorithm visualized (Lanham, M. (2018)).

2.3 Multilayer preceptron (MLP)

A multilayer preceptron is a feedforward artificial neural network (ANN) (Hjort-Jensen, M. (2019)). A preceptron is a single neuron model that contributes to a larger neural network. The term *multilayer* comes from having three or more layers in the network; an **input layer**, one or more **hidden layers** and an **output layer**.




Figure 2: Visualizing of a neural network (Hassan, M.H. et al. (2015)).

2.3.1 Neurons

Artificial neurons are the building blocks of the neural network. These artificial neurons have the same mission as biological neurons, to receive a weighted input signal $\mathbf{x}_i \mathbf{w}_i$ and produce an output signal \mathbf{y} by using the activation function σ .




Figure 3: Artificial neuron (Chandra, L.A. (2012)).

2.3.2 Activation function

When the neuron receive the weighted inputs, the signal will pass through and reach the activation function. The function will map the weighted sum from the input to the output of the neuron. As mentioned, the activation function is a non-linear function. In this report the sigmoid function and ReLU were used. The mathematical model for the output signal is given as

$$y = f\left(\sum_{i=1}^n w_i x_i + b_i\right) = f(z) \quad (5)$$

The activation function receives the input signal x_i . The activation is defined as $z = (\sum_{i=1}^n w_i x_i + b_i)$. The output y here is the input x_i for the next layer. This connection does the multilayer preceptron to a fully-connected neural network.

To calculate the weighted sum we can take a look at the mathematical model for weighted sum in the first hidden layer

$$z_i^1 = \sum w_{ij}^1 x_i = b_i^1 \quad (6)$$

And with a more general formal for the activation function, where N_ℓ is the

number of neurons in layer ℓ .

$$y_i^\ell = f^\ell(z_i^\ell) = f^\ell\left(\sum_{j=1}^{N_{\ell-1}} w_{ij}^\ell y_j^{\ell-1} + b_i^\ell\right) \quad (7)$$

There are many activation functions in Machine Learning, but as mentioned, the Sigmoid function is the common function for logistic regression (section 2.1.1). The hyperbolic function and ReLU are also commonly used to train neural networks. The hyperbolic function is an alternative to the Sigmoid function and has a feature that lets the network learn even when its initial weights are small values (Sussillo, d. & Abbot, L.F. (2014)) because $\tanh'(x)$ is continuous around $x = 0$. While for Sigmoid and ReLU this can be problematic. The hyperbolic function

$$f(x) = \tanh(x) = \frac{2}{1 + \exp^{-2x}} - 1 \quad (8)$$

ReLU is short for Rectified Linear Unit and is the most common activation function in deep learning. ReLU is a non-linear activation function and are more efficient for training large networks (Yang, J. (2017)). The formula for ReLU are

$$f(x) = \begin{cases} 0, & \text{for } x < 0 \\ x, & \text{for } x \geq 0 \end{cases} \quad (9)$$

The following image shows how the various activation functions and their graphs




Figure 4: Activation functions and their graph (Jadon, S. (2018)).

2.3.3 Feedforward and backpropagation

Now that we know how a neuron behaves and what an activation function does, we can explain the processes in a neural network; feedforward and backpropagation. Feedforward calculates the predicted output \hat{y} for each neuron and backpropagation updates the weights and biases (Loy, J. (2018)). The following graph illustrates the processes




Figure 5: Processes in a neural network (Loy, J. (2018)).

As illustrated in figure 5, the weighted input with bias $W_1x + b_1$ are coming in to the input layer. From the input layer, the signal goes through to the hidden layers and activate the activation functions $z = \sigma(W_1x + b_1)$ for every neuron (section 2.3.2) in the hidden layers. We get the total output signal in the last layer $\hat{y} = \sigma(W_2z + b_2)$ and evaluate the model with the

cost function $\hat{\beta}$ (**Loss**(\hat{y}, y) in figure 5).

When we measured the error with the costfunction $C(\hat{\beta})$ (section 2.1), we can now start to the process backpropagation. In order to adjust a reasonable value for the weights and bias. We use the the Stochastic Gradient Methods (section 2.2) to adjust the bias and weights.

2.4 One hot encoding

Working with real data can be complicated if, for instance, the data is categorial data with variables that contain labels rater than numeric values. The numeric values might also not have a relationship between them. This is the case for the dataset we used in this project. By applying one hot encoding, the integer values representing a category for each featuer is being replaced by unique binary labels (0 and 1) for each of these values (see Implementation chapter).

2.5 Model evaluation

A confusion matrix is a table showing us how accurate our classification model is for the given dataset where the true values are known and gives us an insight into the errors (Brownlee, J. (2016)).

Confusion Matrix		Modeled Values: x_m	
		False	True
Actual Values: x	False	True Negatives	False Positives
	True	False Negatives	True Positives

Figure 6: Confusion matrix (Namari, H. & Shaffer, J. (2017)).

Definitions of the terms

- True Negatives is the number of observations who is correctly predicted negative.
- True Positives is the number of observations who is correctly predicted positive.
- False Positives is the number of observed negatives who is predicted positive.
- False Negatives is the number of observed positive, but is predicted negative.

Accuracy is defined as the ratio of true values who are correctly assigned by the model to the total number of observed values (Brownlee, J. (2016)). Accuracy can be defined by the equation following

$$\text{Accuracy} = \frac{\text{True Positives} + \text{True Negatives}}{\text{total number of observations}}$$

The goal is to have a high accuracy, where the value 1 is the best.

The F1-score (or just F-score) is a measure of a test's accuracy, we need precision p and the recall r of the test to compute the F-score. The function reaches its best value at 1 and worst at 0.

Recall are defined as the ratio of correctly classified positive outcomes divides by the total number of positives (Brownlee, J. (2016)). High value of recall indicates small number of False Negative, that means the dataset is correctly classified.

$$\text{Recall} = \frac{\text{True Positive}}{\text{True Positive} + \text{False Negative}}$$

Precision are defined as the ratio of correctly classified positive outcomes divided by the total number of originally positive input. High precision

shows that positive outcomes is indeed positive, that means small number of False Positive.

$$\text{Precision} = \frac{\text{True Positive}}{\text{True Positive} + \text{False Positive}}$$

The F_1 -score is the harmonic mean of the precision and recall

$$F_1 = 2 \cdot \frac{\text{recall} \cdot \text{precision}}{\text{recall} + \text{precision}}$$

R2-score measure the variance explained by the model divided by the total variance, and tell us how well the regression model fits the dataset by percentage from a scale 0 to 100 % (Tangen, M. & Le, T. (2019)).

$$R^2(\hat{y}, \tilde{y}) = 1 - \frac{\sum_{i=0}^{n-1} (y_i - \tilde{y}_i)^2}{\sum_{i=0}^{n-1} (y_i - \bar{y})^2},$$

where \bar{y} is the mean of the values.

Mean Squared Error (MSE) taking the average of the square of the values, where the values is the differences between the observed values y and the predicted ones \hat{y} . The bigger the values ii the larger the error and 0 means that the regression model is perfect (Tangen, M. & Le, T. (2019)).

$$MSE(\hat{y}, \tilde{y}) = \frac{1}{n} \sum_{i=0}^{n-1} (y_i - \tilde{y}_i)^2,$$

3 Data

3.1 Credit card data

This dataset is retrieved from the Machine Learning Repository of the University of California, School of Information and Computer Science (Yeh, I-C. & Lien, C-H. (2009)). It is a collection of payment data from a bank i

Taiwan in 2005.

The dataset has 23 explanatory variables ($X_1 - X_{23}$) in addition to client IDs and the response variable (Y). The explanatory variables holds information on the different consumers, such as information on their gender, age and education. The Appendix gives a detailed description of alle the explanatory variables. The response variable Y describes whether or not the client defaults in the next month, i.e. if the client fails to pay by the deadline or not. Y can have the value 1 (meaning that the client did fail to pay) or 0 (meaning that the client payed on time).

3.2 Franke function

The two-dimensional Franke function is a weighted sum of four exponentials. This function is commonly used for testing various fitting algorithms like we did in project 1 (Tangen, M. & Le, T. (2019)). The function is following

$$f(x, y) = \frac{3}{4} \exp \left(-\frac{(9x-2)^2}{4} - \frac{(9y-2)^2}{4} \right) + \frac{3}{4} \exp \left(-\frac{(9x+1)^2}{49} - \frac{(9y+1)}{10} \right) + \frac{1}{2} \exp \left(-\frac{(9x-7)^2}{4} - \frac{(9y-3)^2}{4} \right) - \frac{1}{5} \exp \left(-(9x-4)^2 - (9y-7)^2 \right) \quad (10)$$

4 Implementation

This project had three parts to implement: logistic regression, artifical neural network for classification and artificial neural network for linear regression. The artifical neural network we developed without using scikit-learn (and similair libraries) are reffered to sometimes in the results as "NN" and does at this point only have the option of 1 hidden layer.

4.1 Data preprocessing

As mentioned above, one should onehot encode the data when it represents different classes. We onehot-encoded columns 2, 3, 4, and 5, i.e. the gender,

education, marital status and age columns. That is because the values in these columns do not have a set relationship between them. All columns holding information on payments always represent a month in between the different values, whereas the columns 2, 3, 4 and 5 does not have such a "set relationship"; there is not a month or another set number/meaning between graduate school and university, and university and high school for instance.

In addition, we removed invalid values and scaled the data (using scikit learn). We then used scikit-learn *train_test_split* to split the data into training and test data. This function shuffles the data randomly before splitting. We ran our analyses using training and test sets of equals sizes, i.e. we used half of the data for training and the other half for testing. As our network performed poorly, we tried to increase the training amount, but this did not change the performance, as shown in figure 7 (compare with figures 10 and 13, which are the performances of our implementation and scikit-learns MLPClassifier for a set η and λ , run with 50/50 training and test data). It would have been interesting to create a more balanced data set, simply by taking out a number of 0s to make the distribution of 0s and 1s 50/50.

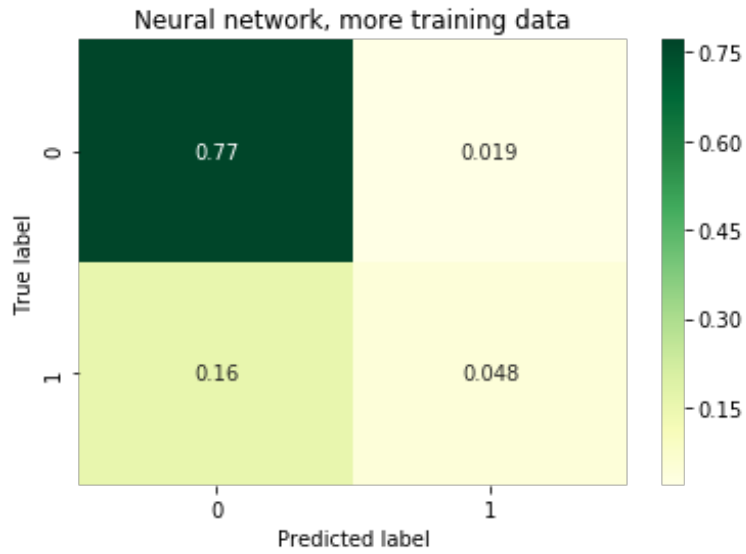


Figure 7: Confusion matrix, more training data.

4.2 Evaluation metrics

As ways of evaluating the performance of our network, we implemented the accuracy, F1, MSE and R2 scores using the scikit-learn library. In addition, we implemented the confusion matrix.

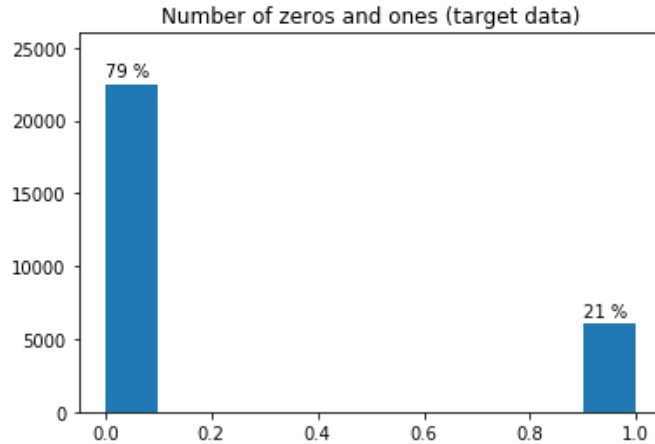


Figure 8: Distribution of 0s and 1s in target data.

As shown in the figure 8, the targets of our dataset contains 79 % 0s and 21 % 1s. This means that even if our accuracy score is at 79 %, the network might only have predicted only zeros (see figure 9). Because of this imbalance in the dataset, the F1 score and confusion matrix are better at determining whether our network performs well or not. Early in the process of building our network we got decent accuracy scores, around ~ 0.78 . However, when plotting the confusion matrix and F1 score, we discovered that the network only predicted one of the classes (0). The decent-looking accuracy score results from the fact that our dataset is very imbalanced. In this analysis, we have therefore focused mainly on the F1 score and confusion matrix to evaluate the network performance, as this takes imbalance into account.

We also applied a regularization parameter to our network, to make sure that the weights do not grow out of control and overfit the training data.

The results from our own implemented neural network were compared to those of scikit-learns MLPClassifier.

We used the MSE and R2 scores to evaluate how the network did on linear regression. Since linear regression is not a binary classification problem, we had to change our model to output actual z values, rather than binary

values.

5 Results & discussion

5.1 Neural network

As mentioned above, our network had issues of only classifying one category. To test if the network was actually learning, we trained it on a dataset only containing 0s and only 1s. We can see in figure 9 that the network was able to learn to classify both, which confirmed that the network was able to learn. To improve performance, we then wanted to change the different hyperparameters.




Figure 9: Confusion matrix, only trained on 0s.




Figure 10: Confusion matrix for neural network




Figure 11: F1 score (MLPClassifier) as a function of η (learning rate) and λ (regularization)




Figure 12: Accuracy score (MLPClassifier) as a function of η (learning rate) and λ (regularization)

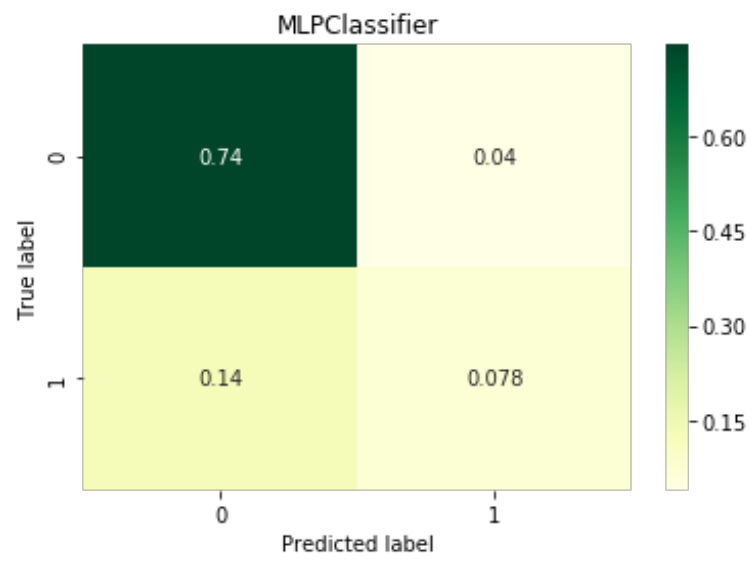


Figure 13: Confusion matrix for MLPClassifier




Figure 14: Heatmap for training data (accuracy score)




Figure 15: Heatmap for test data (accuracy score)

Figures 14 and 15 show how the accuracy score of our neural network changes with different sets of learning rates and regularization parameters, for training and test data respectively. We see little difference between the two cases. We identify an ideal learning rate η of 0.01-0.1.




Figure 16: Heatmap for test data (F1 score)




Figure 17: Heatmap for test data (F1 score)

Figures 16 and 17 show how the F1 score changes with different learning rates and regularization parameters.

In general, we see that the learning rate influences the accuracy more than the regularization parameter.

5.2 Franke's function

Figures 18 and 19 shows the MSE score on the Franke function as a function of lambda and eta. As we can see, our implementation is performing much worse than that of scikit-learn, which might point to us having failed in changing the network to run for linear regression. There are, however, some combinations of lambdas and etas that yield good values. The same is evident

when looking at the R2 scores (figures 20 and 21).




Figure 18: Linear regression (scikit-learn), MSE as a function of η (learning rate) and λ (regularization)




Figure 19: Linear regression (NN), MSE as a function of η (learning rate) and λ (regularization)




Figure 20: Linear regression (scikit-learn), R2 as a function of η (learning rate) and λ (regularization)

Metric	OLS	Ridge	LASSO	Neural network	MLPRegressor
R^2	0.9699	0.9699	0.628	0.74	0.99
MSE	0.0025	0.0090	0.0307	0.0484	0.0023

Table 1: Regression metrics




Figure 21: Linear regression (NN), R^2 as a function of η (learning rate) and λ (regularization)

Table 1 shows the MSE and R^2 scores for the different regression methods Tangen, M. & Le, T. (2019), and results produced in this project. As we can see, the MLPRegressor performs better than all other methods for regression.

6 Conclusion

Our results agree with those from Yeh, I-C. & Lien, C-H. (2009) for the classification case. We see a huge difference in performance of logistic regression and using neural networks. For the regression case, we see that the neural network from scikit-learn beats the other regression methods.

In the future, it would be interesting to try to improve our own neural network code, especially for the regression case, as it performed very badly. We suspect that there is an error we have yet to spot. From there, it would be interesting to see if we could make the network more flexible in terms of number of layers. It would also have been very interesting to compare the results if we were to make a dataset that was not biased, i.e. one containing the same amount of 0s and 1s in the target data. Our guess is that this would have helped the performance. Furthermore, we could have looked at different ways of preprocessing the data, in terms of one hot encoding and scaling of the data, which also could have helped performance. We realize that preprocessing the data properly is one of the most (if not the most) important part of projects and analysis like this, and that spending too little time on it can lead to problems in the algorithms. We will to keep this in the back of our heads going into future projects and treat real, messy data with greater care.

7 References

Brownlee, J. (2016) *What is a Confusion Matrix in Machine Learning*.

Downloaded from: <https://machinelearningmastery.com/confusion-matrix-machine-learning/>

Chandra, L.A. (2012) *Perceptron: The Artificial Neuron (An Essential Upgrade To The McCulloch-Pitts Neuron*

Downloaded from: <https://towardsdatascience.com/perceptron-the-artificial-neuron-4d8c70d5cc8d>

Devore, J. L. & Berk, K. N. (2012) *Modern Mathematical Statistics with Applications* 2nd. ed. Springer.

Hassan, M.H. et al. (2015) *Assessment of Artificial Neural Network for bathymetry estimation using High Resolution Satellite imagery in Shallow Lakes: Case Study El Burullus Lake*. Conference: Eighteenth International Water Technology Conference (2015).

Hjort-Jensen, M. (2019). *Data Analysis and Machine Learning: Neural networks, from the simple perceptron to deep learning*. Lecture notes, FYS-STK4155 Applied data analysis and machine learning. University of Oslo.

Hjort-Jensen, M. (2019). *Data Analysis and Machine Learning: Optimization and Gradient Methods*. Lecture notes, FYS-STK4155 Applied data analysis and machine learning. University of Oslo.

Hjort-Jensen, M. (2019). *Data Analysis and Machine Learning: Logistic Regression*. Lecture notes, FYS-STK4155 Applied data analysis and machine learning. University of Oslo.

Jadon, S. (2018). *Introduction on Different Activation Functions for Deep Learning*. Downloaded from: <https://medium.com/@shrutijadon10104776/survey->

on-activation-functions-for-deep-learning-9689331ba092

Lanham, M. (2018). *Learn ARCore - Fundamentals of Google ARCore*. Publisher: Packt Publishing.

Loy, J. (2018). *How to build your own Neural Network from scratch in Python*.

Downloaded from: <https://towardsdatascience.com/how-to-build-your-own-neural-network-from-scratch-in-python-68998a08e4f6>

Namari, H. & Shaffer, J. (2017). *Visualizing a Confusion Matrix*. Downloaded from: <http://www.dataplusscience.com/VizConfusion.html>

Sussillo, d. & Abbot, L.F. (2014). *Random Walk Initialization for Training Very Deep Feedforward Networks*. Conference: The International Conference on Learning Representations ICLR (2015).

Tangen, M. & Le, T. (2019) *Project1 FYS-STK4155*. University of Oslo, Oslo.

Yang, J. (2017). *ReLu and Softmax Activation Functions*. Downloaded from: <https://github.com/Kulbear/deep-learning-nano-foundation/wiki/ReLU-and-Softmax-Activation-Functions>

Yeh, I-C. & Lien, C-H. (2009) *The comparisons of data mining techniques for the predictive accuracy of probability of default of credit card clients*. Expert Systems with Applications, 36(2), 2473-2480.

8 Appendix

8.1 Explanatory variables

Variable	Description	Unit / value range	Name
X_1	Given credit	NT dollars	LIMIT_BAL
X_2	Gender	1 - 2	SEX
X_3	Education	1 - 4	EDUCATION
X_4	Marital status	1 - 3	MARRIAGE
X_5	Age	Year	AGE
$X_6 - X_{11}$	History of past pay	-1, 1 - 9	PAY_1 ¹ - PAY_6
$X_{12} - X_{17}$	Amount of bill statement	NT dollars	BILL_AMT1 - BILL_AMT6
$X_{18} - X_{23}$	Amount of previous payment	NT dollars	PAY_AMT1 - PAY_AMT6

Table 2: Overview of explanatory variables of the Credit card data set

Variable X_1 holds the amount of credit (given in NT dollars), and includes both the individual and supplementary credit (i.e., the consumers own credit and the credit of the family of that consumer).

Variable X_2 describes the gender, and has a value of 1 if the consumer is male and 2 if the consumer is female.

Variable X_3 describes the consumer's level of education, and takes on values from 1 to 4, where 1 is graduate school, 2 university, 3 high school and 4 others.

Variable X_4 describes whether the consumer is married ($X_4 = 1$), single ($X_4 = 2$) or some other marital status ($X_4 = 3$).

Variable X_5 holds the age of the consumer in years.

Variables $X_6 - X_{11}$ holds past monthly payment records, from April (X_{11}) to September (X_6) 2005. These variables take on values from 1-9, which describes the delay of the past payment in months. E.g., if one of the variables has value 1, then the payment was done one month after the deadline and 9 months (or more) if the variables has the value 9. If the payment was done on time, these variables has value -1.

Variables $X_{12} - X_{17}$ holds the amount on the bill statement (given in NT dollars) for the same months as for the previous variables, whereas variables $X_{18} - X_{23}$ holds the amount of the previous payments (given in NT dollars), also for those same months. As for the past monthly payment records, the variables are ordered backwards (starting in September (X_{12}/X_{18}) going to April (X_{17}/X_{23})).